

CAFFE-DEEP LEARNING FRAMEWORK

Omar Hassaan
LUMS MS(CS)

Research Associate @ITU
omar.hassaan@gmail.com

LAYOUT

1

- Brief Introduction
- Basic requirements for setting up Caffe Environment

2

- Download and Install Caffe
- Train GoogLeNet on MNIST Data set

3

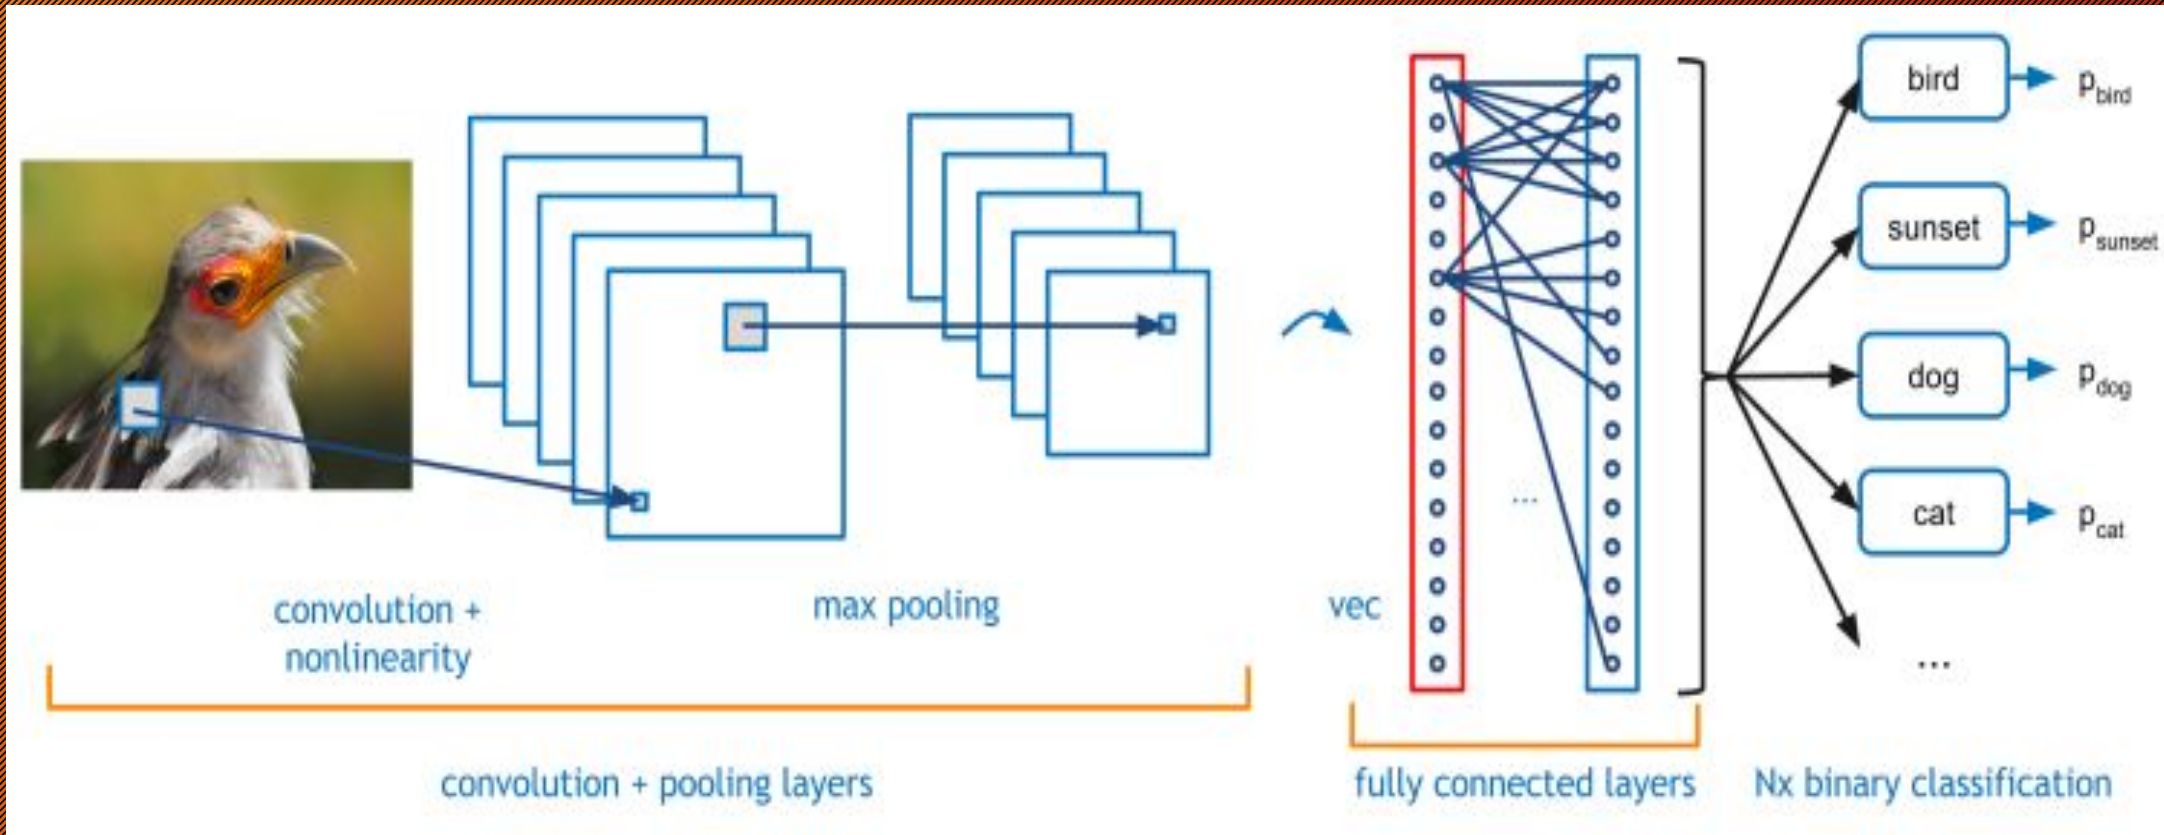
- Testing a trained model with test image
- Training your own dataset

4

- Extract Features from any layer in Caffe

What is Caffe?

- Caffe = Convolution Architecture for Fast Feature Embedding



What is Caffe?

- Open framework for Deep learning.
- Pure C++ / CUDA architecture for deep learning
- Command line, Python, MATLAB interfaces
- Fast, well-tested code
- Tools, reference models, demos, and recipes
- Seamless switch between CPU and GPU
- Offers model definitions, optimization settings and pretrained weights.

Setting Up Caffe Environment

- Core2duo/i3/i5/i7 machine
- Sufficient amount of RAM
- GPU- shared/dedicated. (Caffe can also work without GPU)
- Dual boot Linux distribution. (NO virtual machine !)
- Internet connection
- Patience !

Installing Caffe

- Setup CUDA
 - `http://developer.download.nvidia.com/compute/cuda/6_5/rel/installers/cuda_6.5.14_linux_64.run`
 - `chmod +x cuda_6.5.14_linux_64.run`
 - `sudo ./cuda_6.5.14_linux_64.run --kernel-source-path=/usr/src/linux-headers-`uname -r`/`
 - Accept the EULA
 - Install the graphics card drivers. (If GPU)
 - Install the toolkit (leave path at default)
 - Install symbolic link
 - Install samples (leave path at default)
- Update Library Path
 - `cp /etc/bash.bashrc /home/user/.bashrc`
 - `echo 'export PATH=/usr/local/cuda/bin:$PATH' >> ~/.bashrc`
 - `echo 'export LD_LIBRARY_PATH=$LD_LIBRARY_PATH:/usr/local/cuda/lib64:/usr/local/lib' >> ~/.bashrc`
 - `source ~/.bashrc`

Download Caffe

- Download Caffe
 - <https://github.com/omair18/Caffe-Setup-Testing-Training>
- Extract **caffe-master.zip** in your Home directory (/home/username/)

Install Dependencies

- Install dependencies
 - `cd caffe-master` (I'll call this folder as \$CAFFE-ROOT)
 - `./linux_dep.sh`
 - `./python_dep.sh`
 - Patience ! :)
- Add a couple of Symbolic links for some reason
 - `sudo ln -s /usr/include/python2.7/ /usr/local/include/python2.7`
 - `sudo ln -s /usr/local/lib/python2.7/dist-packages/numpy/core/include/numpy/ <space> /usr/local/include/python2.7/numpy`

Compiling Caffe

- Compile Caffe
 - `cp Makefile.config.example Makefile.config`
 - `nano Makefile.config`
 - If you don't have GPU, uncomment the line `#CPU_ONLY := 1`
 - Under `PYTHON_INCLUDE` replace `/usr/lib/python2.7/dist-packages/numpy/core/include` with `/usr/local/lib/python2.7/dist-packages/numpy/core/include` (i.e. add /local)
- Start Compilation
 - `make pycaffe` //for Python
 - Optional:
 - `Make matcaffe` //for matlab
 - `make all`
 - `make test`
 - Patience :)

Testing Installation

- Download the ImageNet Caffe model and labels
 - `./scripts/download_model_binary.py models/bvlc_reference_caffenet`
 - `./data/ilsvrc12/get_ilsvrc_aux.sh`
- First Download ilsvrc12 package
 - `$CAFFE_ROOT/data/get_ilsvrc_aux.sh`
 - `cp $CAFFE_ROOT/caffe_ilsvrc12/synset_words.txt $CAFFE_ROOT/data/ilsvrc12/`
- Test your installation by running the ImageNet model on an image of a kitten
 - `python python/classify.py --print_results examples/images/cat.jpg foo`
 - Expected result: `[('tabby', '0.27933'), ('tiger cat', '0.21915'), ('Egyptian cat', '0.16064'), ('lynx', '0.12844'), ('kit fox', '0.05155')]`

1

- Brief Introduction
- Basic requirements for setting up Caffe Environment

2

- Download and Install Caffe
- Train GoogLeNet on MNIST Data set

3

- Testing a Trained model
- Training your own dataset

- Extract Features from any layer in Caffe

Training LeNet on MNIST dataset

- Change directory to \$CAFFE_ROOT
- Prepare dataset
 - `./data/mnist/get_mnist.sh`
 - `./examples/mnist/create_mnist.sh`
- We will get 2 lmdb files from the above step
 - `Mnist_train_lmdb`
 - `Mnist_test_lmdb`
- All training layers are written in *.prototxt files
 - `Data/Convolution/Pooling/Inner product/ReLU/Loss`
- Define Mnist solver file for training/test protocol buffer definition
 - `$CAFFE_ROOT/examples/mnist/lenet_solver.prototxt`
 - Contains training iterations, base learning rate, max # iterations and solver mode

Network Layers(Vision Layers)

- Vision layers usually take images as input and produce images as output.
- Most of the vision layers work by applying a particular operation to some region of the input to produce a corresponding region of the output
- Examples
 - Convolution Layer
 - Pooling Layer
 - Max layer ..

Convolution Layer

- Layer type: Convolution
- CPU implementation:
 - `$CAFFE_ROOT/src/caffe/layers/convolution_layer.cpp`
- Parameters (ConvolutionParameter convolution_param)
- num_output (c_o):
 - the number of filters
- kernel_size :
 - specifies height and width of each filter
- weight_filler [default type: 'constant' value: 0]
- bias_term [default true]:
 - specifies whether to learn and apply a set of additive biases to the filter outputs
- pad [default 0]:
 - specifies the number of pixels to (implicitly) add to each side of the input
- stride [default 1]:
 - specifies the intervals at which to apply the filters to the input

```
layer {
  name: "conv1"
  type: "Convolution"
  bottom: "data"
  top: "conv1"
  # learning rate and decay multipliers for the filters
  param { lr_mult: 1 decay_mult: 1 }
  # learning rate and decay multipliers for the biases
  param { lr_mult: 2 decay_mult: 0 }
  convolution_param {
    num_output: 96      # learn 96 filters
    kernel_size: 11     # each filter is 11x11
    stride: 4           # step 4 pixels between each filter application
    weight_filler {
      type: "gaussian" # initialize the filters from a Gaussian
      std: 0.01        # distribution with stdev 0.01 (default mean: 0)
    }
    bias_filler {
      type: "constant" # initialize the biases to zero (0)
      value: 0
    }
  }
}
```


Pooling Layer

- Layer type: Pooling
- CPU implementation:
 - `$CAFFE_ROOT/src/caffe/layers/pooling_layer.cpp`
- Parameters (PoolingParameter pooling_param)
- kernel_size (or kernel_h and kernel_w):
 - specifies height and width of each filter
- pool [default MAX]:
 - the pooling method. Currently MAX, AVE, or STOCHASTIC
- pad (or pad_h and pad_w) [default 0]: specifies the number of pixels to (implicitly) add to each side of the input
- stride (or stride_h and stride_w) [default 1]:
 - specifies the intervals at which to apply the filters to the input
- pad (or pad_h and pad_w) [default 0]:
 - specifies the number of pixels to (implicitly) add to each side of the input
- stride (or stride_h and stride_w) [default 1]:
 - specifies the intervals at which to apply the filters to the input

```
layer {
  name: "pool1"
  type: "Pooling"
  bottom: "conv1"
  top: "pool1"
  pooling_param {
    pool: MAX
    kernel_size: 3 # pool over a 3x3 region
    stride: 2      # step two pixels (in the bottom blob) between pooling regions
  }
}
```


Loss Layer

- Loss drives learning by comparing an output to a target and assigning cost to minimize
- Softmax:
 - Layer-type: SoftmaxWithLoss
- Sum of Squares/Euclidean
 - Layer-type: EuclideanLoss
- Hinge/Margin
 - Layer type: HingeLoss

```
# L1 Norm
layer {
  name: "loss"
  type: "HingeLoss"
  bottom: "pred"
  bottom: "label"
}

# L2 Norm
layer {
  name: "loss"
  type: "HingeLoss"
  bottom: "pred"
  bottom: "label"
  top: "loss"
  hinge_loss_param {
    norm: L2
  }
}
```


ReLu Layer (Neuron Layer)

- Neuron layers are element-wise operators, taking one bottom blob and producing one top blob of the same size.
- Layer type: ReLU
 - Sigmoid, TanH
 - <http://caffe.berkeleyvision.org/tutorial/layer.html>
- CPU implementation:
 - `$CAFFE_ROOT/src/caffe/layers/relu_layer.cpp`
- Parameters (ReLUParameter relu_param)
- negative_slope [default 0]:
 - specifies whether to leak the negative part by multiplying it with the slope value rather than setting it to 0.

```
layer {  
  name: "relu1"  
  type: "ReLU"  
  bottom: "conv1"  
  top: "conv1"  
}
```


Training LeNet on MNIST dataset

- Start training using
 - `$CAFFE_ROOT/examples/mnist/train_lenet.sh`
 - `./build/tools/caffe train --solver=examples/mnist/lenet_solver.prototxt`
- End result will be a *.caffemodel file containing our trained network.

1

- Brief Introduction
- Basic requirements for setting up Caffe Environment

2

- Download and Install Caffe
- Train GoogLeNet on MNIST Data set

3

- **Testing a Trained model**
- Training your own dataset

- Extract Features from any layer in Caffe

Testing a Trained Model

- Create a LMDB database of the testing image set.
- You must have prototxt files available and also the snapshot of trained model.
- Use the following python script to test your trained mode.
 - `Test_your_own_model.py`

1

- Brief Introduction
- Basic requirements for setting up Caffe Environment

2

- Download and Install Caffe
- Train GoogLeNet on MNIST Data set

3

- Testing a trained model
- **Training your own dataset**

4

- Extract Features from any layer in Caffe

Creating LMDB with image dataset

- Create 2 files containing [path/to/image <space> label]
 - Train.txt
 - Val.txt
- Reference example
 - `$CAFFE_ROOT/examples/imagenet/create_imagenet.sh`
- Download ilsvrc12 package by running
 - `$CAFFE_ROOT/data/get_ilsvrc_aux.sh`
 - Open `$CAFFE_ROOT/data/ilsvrc12/train.txt` or `val.txt`
- Run the *convert_imageset* program to convert all the images to lmdb format using
 - `$CAFFE_ROOT/build/tools/convert_imageset [FLAGS] ROOTFOLDER LISTFILE DB_NAME`
 - FLAGS can be gray, shuffle, backend, resize_width, resize_height, check_size, encoded, encoded type.
 - ROOTFOLDER = folder containing all the images
 - LISTFILE = train.txt/val.txt for creating training/testing database

Training with your own Dataset

- After creating your own LMDB database, it's time to create prototxt files containing all the layers information of your network.
 - References
 - `$CAFFE_ROOT/examples/mnist/lenet.prototxt`
 - `$CAFFE_ROOT/examples/mnist/lenet_train_test.prototxt`
 - `$CAFFE_ROOT/examples/mnist/lenet_solver.prototxt`
 - `$CAFFE_ROOT/examples/mnist/train_lenet.sh`
- Change path of prototxt files in `train_lenet.sh` file(see “Training LeNet on MNIST using Caffe” for reference)
- Start training by running `train_lenet.sh`

1

- Brief Introduction
- Basic requirements for setting up Caffe Environment

2

- Download and Install Caffe
- Train GoogLeNet on MNIST Data set

3

- Testing a trained model
- Training your own dataset

4

- Extract Features from any layer in Caffe

Create Mean of Imageset(binaryproto)

- Crop all images to 227x227 or 256x256
- Create LMDB database of the image set .
- Now run
 - `$CAFFE_ROOT/build/tools/compute_image_mean [LMDB] [FILE NAME]`
 - **LMDB** is the database you created
 - **FILE NAME** is the destination filename that will be created with .binaryproto extension

Extract Features from Any layer in Caffe

- Crop all images to 227x227 or 256x256 & Create LMDB database of the image set .
- Create a prototxt file, in which we have to mention our LMDB database and mean of image set (binaryproto)
 - Reference `$CAFFE_ROOT/examples/feature_extraction/imagenet_val.prototxt`
- Now Run
 - `$CAFFE_ROOT/build/tools/extract_features.bin [MODEL File] [NETWORK] [LAYER] [DESTINATION] [BATCHSIZE] [LMDB]`
 - **MODEL FILE** = pretrained caffe model
 - `$CAFFE_ROOT/models/bvlc_reference_caffe_net/bvlc_caffenet.caffemodel`
 - **NETWORK** = Prototxt file we just made in step 3
 - **LAYER** = Layer name whose features we want to extract e.g fc7
 - **DESTINATION** = Folder name to save results
 - **BATCH SIZE** e.g 10
 - **LMDB** = to store results in LMDB format