

CAFFE-DEEP LEARNING FRAMEWORK

Omais Hassaan
LUMS MS(CS)

omais.hassaan@gmail.com

LAYOUT

1

- **Brief Introduction**
- Basic requirements for setting up Caffe Environment

2

- Testing Caffe Installation
- Train LeNet Architecture on MNIST Data set

3

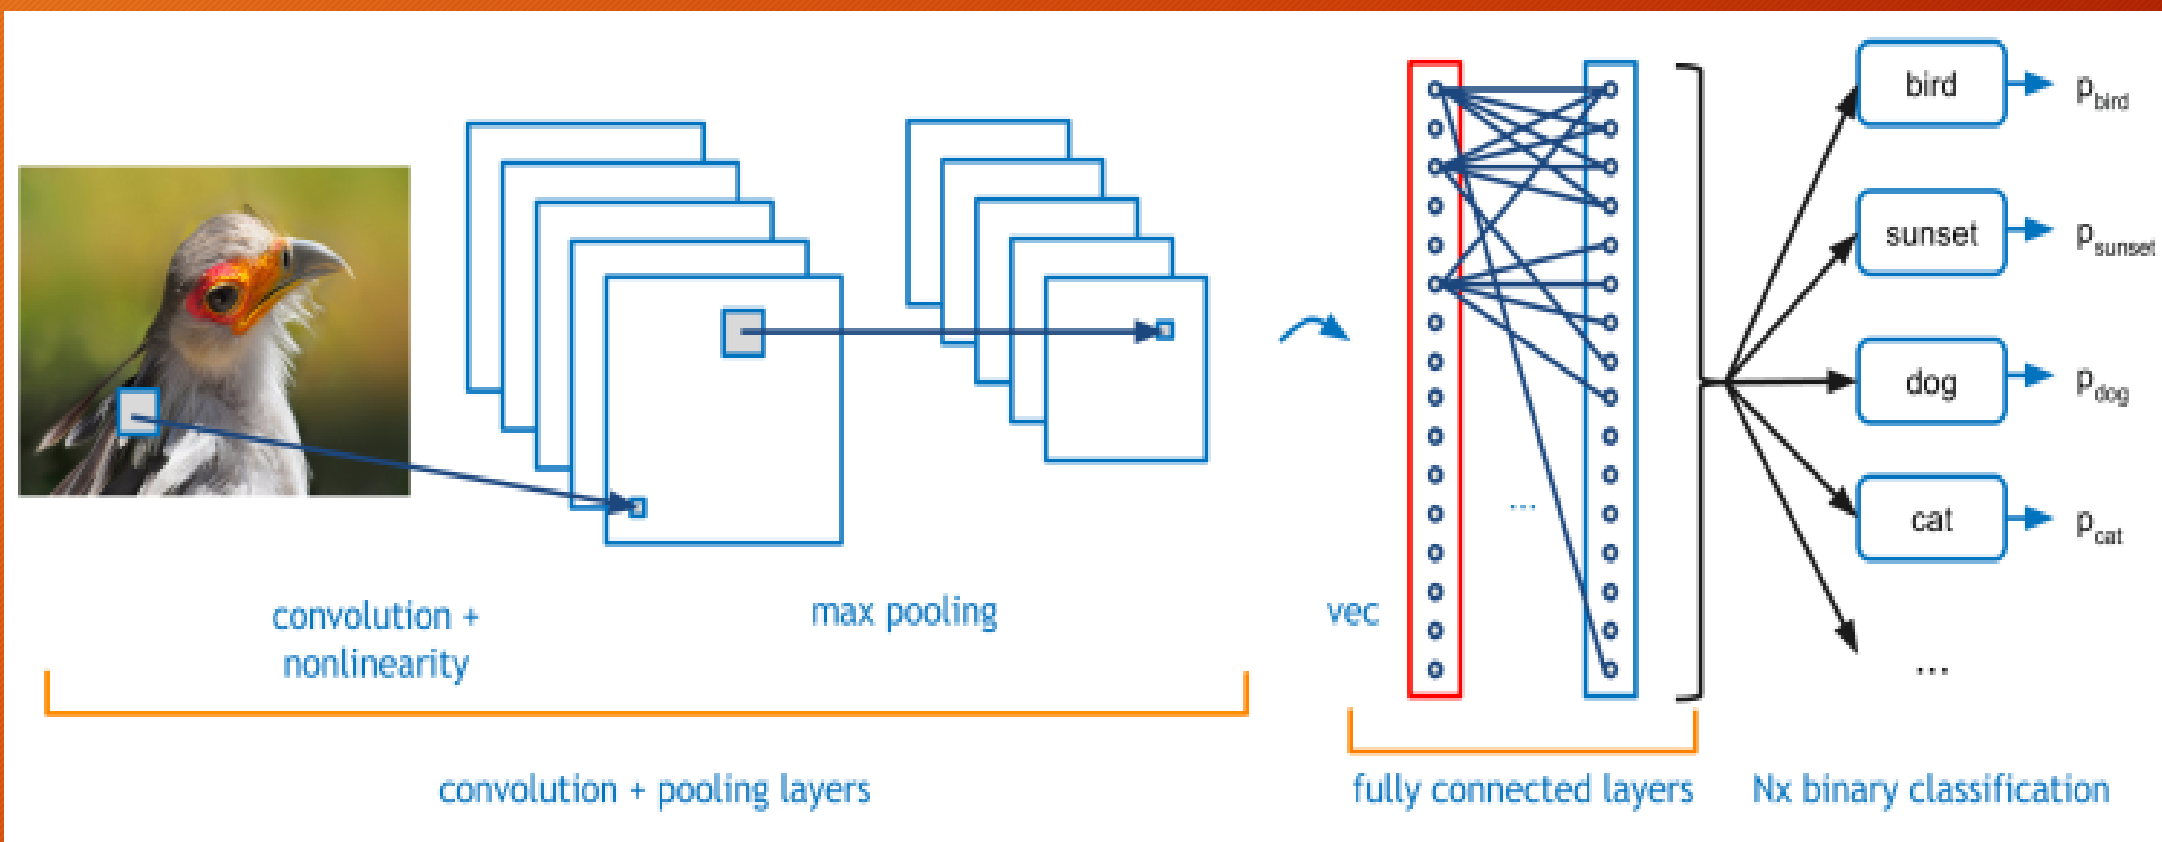
- Training your own Dataset
- Testing a trained Model

4

- Filter Visualization
- Fine Tuning of a Trained Model.

What is Caffe?

- Caffe = Convolution Architecture for Fast Feature Embedding



What is Caffe?

- Open framework for Deep learning.
- Pure C++ / CUDA architecture for deep learning
- Command line, Python, MATLAB interfaces
- Fast, well-tested code
- Tools, reference models, demos, and recipes
- Seamless switch between CPU and GPU
- Offers model definitions, optimization settings and pretrained weights.

LAYOUT

1

- Brief Introduction
- Basic requirements for setting up Caffe Environment

2

- Testing Caffe Installation
- Train LeNet Architecture on MNIST Data set

3

- Training your own Dataset
- Testing a trained Model

4

- Filter Visualization
- Fine Tuning of a Trained Model.

Setting Up Caffe Environment

- Core2duo/i3/i5/i7 machine
- Sufficient amount of RAM
- GPU- shared/dedicated. (Caffe can also work without GPU)
- Dual boot Linux distribution. (recommended)
- Internet connection
- Patience !

LAYOUT

1

- Brief Introduction
- Basic requirements for setting up Caffe Environment

2

- **Testing Caffe Installation**
- Train LeNet Architecture on MNIST Data set

3

- Training your own Dataset
- Testing a trained Model

4

- Fine Tuning of a Trained Model.

Testing Installation

- Download the ImageNet Caffe model and labels
 - `./scripts/download_model_binary.py models/bvlc_reference_caffenet`
 - `./data/ilsvrc12/get_ilsvrc_aux.sh`
- First Download ilsvrc12 package
 - `$CAFFE_ROOT/data/get_ilsvrc_aux.sh`
 - `cp $CAFFE_ROOT/caffe_ilsvrc12/synset_words.txt $CAFFE_ROOT/data/ilsvrc12/`
- Test your installation by running the ImageNet model on an image of a kitten
 - `python python/classify.py --print_results examples/images/cat.jpg foo`
 - Expected result: `[('tabby', '0.27933'), ('tiger cat', '0.21915'), ('Egyptian cat', '0.16064'), ('lynx', '0.12844'), ('kit fox', '0.05155')]`

LAYOUT

1

- Brief Introduction
- Basic requirements for setting up Caffe Environment

2

- Testing Caffe Installation
- Train LeNet Architecture on MNIST Data set

3

- Training your own Dataset
- Testing a trained Model

4

- Filter Visualization
- Fine Tuning of a Trained Model.

Training LeNet on MNIST dataset

- Change directory to \$CAFFE_ROOT (/home/omair/caffe-master)
- Prepare dataset
 - ./data/mnist/get_mnist.sh
 - ./examples/mnist/create_mnist.sh
- We will get 2 lmdb files from the above step
 - mnist_train_lmdb (examples/mnist/)
 - mnist_test_lmdb
- All training layers are written in *.prototxt files
 - Data/Convolution/Pooling/Inner product/ReLU/Loss
- Define Mnist solver file for training/test protocol buffer definition
 - \$CAFFE_ROOT/examples/mnist/lenet_solver.prototxt
 - Contains training iterations, base learning rate, max # iterations and solver mode

Network Layers

- We define our network architecture in “prototxt “ files
- There are 2 kinds of protoxt files.
- TRAIN_TEST.PROTOTXT
 - Define Convolution, Pooling, Softmax layers
- FOO_SOLVER.PROTOTXT
 - Defines Learning parameters of our model.

Network Layers(Vision Layers)

- Vision layers usually take images as input and produce images as output.
- Most of the vision layers work by applying a particular operation to some region of the input to produce a corresponding region of the output
- Examples
 - Convolution Layer
 - Pooling Layer
 - Max layer ..

Convolution Layer

- Layer type: Convolution
- CPU implementation:
 - `$CAFFE_ROOT/src/caffe/layers/convolution_layer.cpp`
- Parameters (ConvolutionParameter convolution_param)
- num_output (c_o):
 - the number of filters
- kernel_size :
 - specifies height and width of each filter
- weight_filler [default type: 'constant' value: 0]
- bias_term [default true]:
 - specifies whether to learn and apply a set of additive biases to the filter outputs
- pad [default 0]:
 - specifies the number of pixels to (implicitly) add to each side of the input
- stride [default 1]:
 - specifies the intervals at which to apply the filters to the input

```
layer {
  name: "conv1"
  type: "Convolution"
  bottom: "data"
  top: "conv1"
  # learning rate and decay multipliers for the filters
  param { lr_mult: 1 decay_mult: 1 }
  # learning rate and decay multipliers for the biases
  param { lr_mult: 2 decay_mult: 0 }
  convolution_param {
    num_output: 96      # learn 96 filters
    kernel_size: 11     # each filter is 11x11
    stride: 4           # step 4 pixels between each filter application
    weight_filler {
      type: "gaussian" # initialize the filters from a Gaussian
      std: 0.01        # distribution with stdev 0.01 (default mean: 0)
    }
    bias_filler {
      type: "constant" # initialize the biases to zero (0)
      value: 0
    }
  }
}
```

Pooling Layer

- Layer type: Pooling
- CPU implementation:
 - `$CAFFE_ROOT/src/caffe/layers/pooling_layer.cpp`
- Parameters (PoolingParameter pooling_param)
- kernel_size (or kernel_h and kernel_w):
 - specifies height and width of each filter
- pool [default MAX]:
 - the pooling method. Currently MAX, AVE, or STOCHASTIC
- pad (or pad_h and pad_w) [default 0]: specifies the number of pixels to (implicitly) add to each side of the input
- stride (or stride_h and stride_w) [default 1]:
 - specifies the intervals at which to apply the filters to the input
- pad (or pad_h and pad_w) [default 0]:
 - specifies the number of pixels to (implicitly) add to each side of the input
- stride (or stride_h and stride_w) [default 1]:
 - specifies the intervals at which to apply the filters to the input

```
layer {
  name: "pool1"
  type: "Pooling"
  bottom: "conv1"
  top: "pool1"
  pooling_param {
    pool: MAX
    kernel_size: 3 # pool over a 3x3 region
    stride: 2      # step two pixels (in the bottom blob) between pooling regions
  }
}
```


Loss Layer

- Loss drives learning by comparing an output to a target and assigning cost to minimize
- Softmax:
 - Layer-type: SoftmaxWithLoss
- Sum of Squares/Euclidean
 - Layer-type: EuclideanLoss
- Hinge/Margin
 - Layer type: HingeLoss

```
# L1 Norm
layer {
  name: "loss"
  type: "HingeLoss"
  bottom: "pred"
  bottom: "label"
}

# L2 Norm
layer {
  name: "loss"
  type: "HingeLoss"
  bottom: "pred"
  bottom: "label"
  top: "loss"
  hinge_loss_param {
    norm: L2
  }
}
```

ReLu Layer (Neuron Layer)

- Neuron layers are element-wise operators, taking one bottom blob and producing one top blob of the same size.
- Layer type: ReLU
 - Sigmoid, TanH
 - <http://caffe.berkeleyvision.org/tutorial/layers/relu.html>
- CPU implementation:
 - `$CAFFE_ROOT/src/caffe/layers/relu_layer.cpp`
- Parameters (ReLUParameter relu_param)
- negative_slope [default 0]:
 - specifies whether to leak the negative part by multiplying it with the slope value rather than setting it to 0.

```
layer {  
  name: "relu1"  
  type: "ReLU"  
  bottom: "conv1"  
  top: "conv1"  
}
```

SOLVER.PROTOTXT

```
base_lr: 0.01      # begin training at a learning rate of 0.01 = 1e-2
lr_policy: "step"  # learning rate policy: drop the learning rate in "steps"
                  # by a factor of gamma every stepsize iterations
gamma: 0.1         # drop the learning rate by a factor of 10
                  # (i.e., multiply it by a factor of gamma = 0.1)
stepsize: 100000   # drop the learning rate every 100K iterations
max_iter: 350000   # train for 350K iterations total
momentum: 0.9
```


Training LeNet on MNIST dataset

- Change directory to \$CAFFE_ROOT (/home/omair/caffe-master)
- Prepare dataset
 - ./data/mnist/get_mnist.sh
 - ./examples/mnist/create_mnist.sh
- We will get 2 lmdb files from the above step
 - mnist_train_lmdb (examples/mnist/)
 - mnist_test_lmdb
- All training layers are written in *.prototxt files
 - Data/Convolution/Pooling/Inner product/ReLU/Loss
- Define Mnist solver file for training/test protocol buffer definition
 - \$CAFFE_ROOT/examples/mnist/lenet_solver.prototxt
 - Contains training iterations, base learning rate, max # iterations and solver mode

Training LeNet on MNIST dataset

- Start training using
 - `$CAFFE_ROOT/examples/mnist/train_lenet.sh`
 - `./build/tools/caffe train --solver=examples/mnist/lenet_solver.prototxt`
- End result will be a `*.caffemodel` file containing our trained network.

1

- Brief Introduction
- Basic requirements for setting up Caffe Environment

2

- Testing Caffe Installation
- Train LeNet Architecture on MNIST Data set

3

- Training your own Dataset
- Testing a trained Model

4

- Filter Visualization
- Fine Tuning of a Trained Model.

Before Starting Training

Creating LMDB with Image dataset

- Create 2 files containing [path/to/image <space> label]
 - Train.txt
 - Val.txt
- Reference example
 - \$CAFFE_ROOT/examples/imagenet/create_imagenet.sh
- Download ilsvrc12 package by running
 - \$CAFFE_ROOT/data/get_ilsvrc_aux.sh
 - Open \$CAFFE_ROOT/data/ilsvrc12/train.txt or val.txt
- Run the *convert_imageset* program to convert all the images to lmdb format using
 - \$CAFFE_ROOT/build/tools/convert_imageset [FLAGS] ROOTFOLDER LISTFILE DB_NAME
 - FLAGS can be gray, shuffle, backend, resize_width, resize_height, check_size, encoded, encoded type.
 - ROOTFOLDER = folder containing all the images
 - LISTFILE = train.txt/val.txt for creating training/testing database

Create Mean of Image Set (binaryproto)

- Convert/Crop all images to 227x227 or 256x256
- Create LMDB of all images
- `CAFFE_ROOT=/home/omair/caffe-master`
- `LMDB=/home/omair/caffe-master/lmdb_creation/ilsvrc12_train_lmdb`
- `OUTPUT=/home/omair/caffe-master/lmdb_creation/meanilsvrc.binaryproto`
- `$CAFFE_ROOT/build/tools/compute_image_mean $LMDB $OUTPUT`

Training with your own Dataset

- After creating your own LMDB database, it's time to create prototxt files containing all the layers information of your network.
 - References
 - `$CAFFE_ROOT/examples/mnist/lenet.prototxt`
 - `$CAFFE_ROOT/examples/mnist/lenet_train_test.prototxt`
 - `$CAFFE_ROOT/examples/mnist/lenet_solver.prototxt`
 - `$CAFFE_ROOT/examples/mnist/train_lenet.sh`
- Change path of prototxt files in `train_lenet.sh` file(see “Training LeNet on MNIST using Caffe” for reference)
- Start training by running `train_lenet.sh`

1

- Brief Introduction
- Basic requirements for setting up Caffe Environment

2

- Testing Caffe Installation
- Train LeNet Architecture on MNIST Data set

3

- Training your own Dataset
- **Testing a trained Model**

4

- Filter Visualization
- Fine Tuning of a Trained Model.

Testing a Trained Model (C++)

- Use C++ API to implement image classification.
- Sample C++ code is present in `$CAFFE_ROOT/examples/cpp_classification/classification.cpp`
- Gets compiled automatically while compiling caffe
- Input arguments
 - Model file (Defining layer architecture)
 - Trained file
 - Mean file (Mean of dataset)
 - Labels (input labels)
 - Test Image.

Testing a Trained Model (C++)

- `CAFFE_ROOT=/home/omair/caffe-master`
- `PROTOTXT=$CAFFE_ROOT/models/bvlc_reference_caffenet/deploy.prototxt`
- `CAFFEMODEL=$CAFFE_ROOT/models/bvlc_reference_caffenet/bvlc_reference_caffenet.caffemodel`
- `MEANBINARYPROTO=$CAFFE_ROOT/data/ilsvrc12/imagenet_mean.binaryproto`
- `LABELS=$CAFFE_ROOT/data/ilsvrc12/synset_words.txt`

- `$CAFFE_ROOT/build/examples/cpp_classification/classification.bin $PROTOTXT $CAFFEMODEL $MEANBINARYPROTO $LABELS $CAFFE_ROOT/examples/images/cat.jpg`

Testing a Trained Model (Python)

- Create a LMDB database of the testing image set.
- You must have prototxt files available and also the snapshot of trained model.
- Make a directory named “test”
- Copy “`test_your_own_model.py`” (located in [repository](#)) in the newly created directory
- Change paths of prototxt files, mean file, trained model.
- Run this code to test your model.
 - `python test_your_own_model.py`

1

- Brief Introduction
- Basic requirements for setting up Caffe Environment

2

- Testing Caffe Installation
- Train LeNet Architecture on MNIST Data set

3


- Training your own Dataset
- Testing a trained Model

4

- **Filter Visualization**
- Fine Tuning of a Trained Model.

Filter Visualization

- We can visualize filters used in any layer using python script.
- Download “filter_visualization.py” from [github respository](#)
- Copy this code to “test” directory.
- Run !
 - `python filter_visualization.py`



1

- Brief Introduction
- Basic requirements for setting up Caffe Environment

2

- Testing Caffe Installation
- Train LeNet Architecture on MNIST Data set

3

- Training your own Dataset
- Testing a trained Model

4

- Filter Visualization
- Fine Tuning of a Trained Model.

Fine Tuning of a Trained Model

- Takes an already trained model, adapts the architecture and resumes training from model weights.
- Instead of training 1,000,000 images + our new dataset from scratch, lets use already trained model of 1,000,000 images and prepare it for our problem.
- Change last layer (Fully connected layer) in our prototxt file.
- Decrease overall learning rate.
- Increase gamma.
- Idea is, that rest of model should change slowly with new data, but new layers should be learnt fast.

Fine Tuning of a Trained Model

- Create lmdb of train/test image dataset.
- Create label files for train/test dataset.
- Edit train_val.prototxt from **bvlc_reference_caffenet problem.**
- All required files are available in **models/finetune_flicker_style.**
- Start fine tuning process
 - **\$CAFFE_ROOT /build/tools/caffe train -solver models/finetune_flicker_style/solver.prototxt -weights models/bvlc_reference_caffe_net/bvlc_reference_caffenet.caffemodel**

```
layer {
  name: "fc8_flickr"
  type: "InnerProduct"
  bottom: "fc7"
  top: "fc8_flickr"
  # lr_mult is set to higher than for other layers, because this layer i
  param {
    lr_mult: 10
    decay_mult: 1
  }
  param {
    lr_mult: 20
    decay_mult: 0
  }
  inner_product_param {
    num_output: 20
    weight_filler {
      type: "gaussian"
      std: 0.01
    }
    bias_filler {
      type: "constant"
      value: 0
    }
  }
}
layer {
  name: "loss"
  type: "SoftmaxWithLoss"
  bottom: "fc8_flickr"
  bottom: "label"
  top: "loss"
}
```